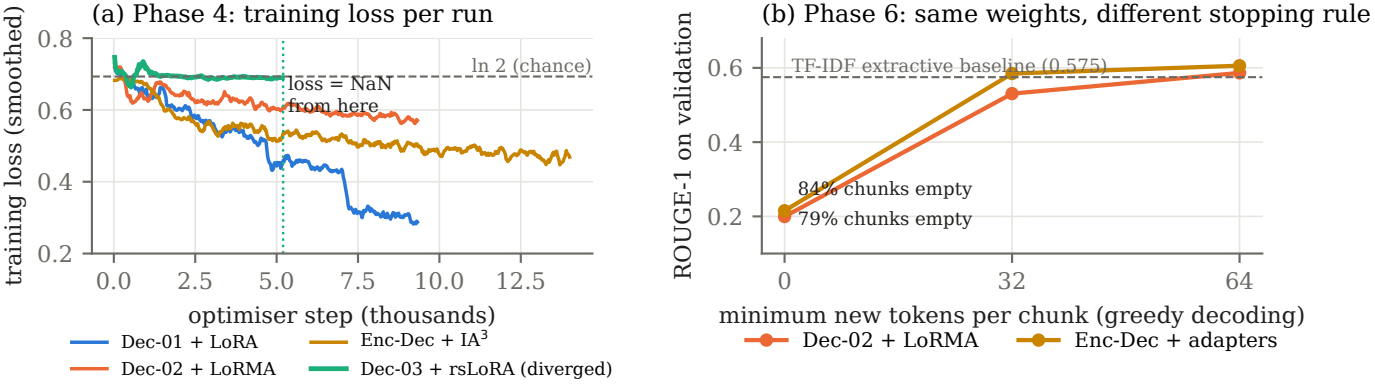
(a) Phase 4: training loss per run
training loss (smoothed)
0.8
0.6
0.4
0.2
ln 2 (chance)
loss = NaN
from here
0.0
2.5
5.0
7.5
10.0
12.5
optimiser step (thousands)
Dec-01 + LoRA
Dec-02 + LoRMA
Enc-Dec + IA³
Dec-03 + rsLoRA (diverged)
(b) Phase 6: same weights, different stopping rule
ROUGE-1 on validation
0.6
0.4
0.2
TF-IDF extractive baseline (0.575)
84% chunks empty
79% chunks empty
0
32
64
minimum new tokens per chunk (greedy decoding)
Dec-02 + LoRMA
Enc-Dec + adapters